

Education

University of Chicago

BA in Mathematics (Honors)

Chicago, IL

September 2010 - June 2014

University of California, Berkeley

Ph.D in Statistics

Berkeley, CA

August 2014 - August 2018

Employment

Vector Institute

Postdoctoral Scholar

Toronto, Canada

August 2018 - Present

- Postdoctoral researcher at the Vector Institute at the University of Toronto
- Advised by Jimmy Ba

Open AI

Research Scientist

San Francisco, USA

March 2016 - November 2017

- Full-time research scientist at Open AI
- Advised by Ilya Sutskever

AI Expert at Y Combinator

Part-time Partner

San Francisco, USA

January 2017 - January 2018

- Advised Y Combinator on the latest trends in AI
- Advised over 50 startups on AI and ML problems

Young Scholars Program/SESAME

Counselor

Chicago, USA

June 2011 - May 2013

- Taught mathematics to Chicagoland high school students and CPS teachers
- Full-time during summers. Every Saturday and Sunday during school year
- Was head counselor from June 2012 - May 2013

Publications

1. *Incentivizing Exploration in Reinforcement Learning with Deep Predictive Models*

Bradly C. Stadie, Sergey Levine, Pieter Abbeel

In Neural Information Processing Systems (NIPS) Deep RL Workshop, Montreal, Canada, December 2015

arXiv: <https://arxiv.org/pdf/1507.00814.pdf>

2. *Third-Person Imitation Learning*

Bradly C. Stadie, Pieter Abbeel, Ilya Sutskever

In the proceedings of the International Conference on Learning Representations (ICLR), Toulon, France, April 2017

arXiv: <https://arxiv.org/pdf/1703.01703.pdf>

3. *One-Shot Imitation Learning*

Yan (Rocky) Duan, Marcin Andrychowicz, Bradly C. Stadie, Jonathan Ho, Jonas Schneider, Ilya Sutskever, Pieter Abbeel, Wojciech Zaremba

In Neural Information Processing Systems (NIPS), Long Beach, California, December 2017

arXiv: <https://arxiv.org/pdf/1703.07326.pdf>

4. *Simulating the Stochastic Dynamics and Cascade Failure of Power Networks*

Charles Matthews, Bradley C. Stadie, Jonathan Weare, Mihai Anitescu, Christopher Demarco

Submitted to IEEE Transactions on Power Systems

arXiv: <https://arxiv.org/pdf/1806.02420.pdf>

5. *Some Considerations on Learning to Explore via Meta-Reinforcement Learning*

Bradley C. Stadie, Ge Yang, Rein Houthooft, Xi Chen, Yan Duan, Yuhuai Wu, Pieter Abbeel, Ilya Sutskever

In Neural Information Processing Systems (NIPS), Montreal, Canada, December 2018

arXiv: <https://arxiv.org/pdf/1803.01118.pdf>

6. *Evolved Policy Gradients*

Rein Houthooft, Richard Y. Chen, Phillip Isola, Bradley C. Stadie, Filip Wolski, Jonathan Ho, Pieter Abbeel

In Neural Information Processing Systems (NIPS) [Spotlight], Montreal, Canada, December 2018

arXiv: <https://arxiv.org/pdf/1802.04821.pdf>

7. *Transfer Learning for Estimating Causal Effects Using Neural Networks.*

Bradley C. Stadie, Soeren R. Kuenzel, Nikita Vemuri, Varsha Ramakrishnan, Jasjeet S. Sekhon, Pieter Abbeel

INFORMS Annual Meeting ML and causal inference workshop

arXiv: <https://arxiv.org/pdf/1808.07804.pdf>

8. *Sampling Aware Reinforcement Learning*

Lunjun Zhang, Bradley C. Stadie, Jimmy Ba

Preprint. Available at: <https://github.com/bstadie/All-Bradly-Stadie-Papers/>

9. *One-Shot Pruning of Recurrent Neural Networks by Jacobian Spectrum Evaluation*

Matthew Zhang, Bradley C. Stadie

Submitted to ICLR 2020

Available at: <https://openreview.net/forum?id=r1e9GCNKvH>

10. *Learning to Learn from Flawed, Failed, and Figurative Demonstrations*

Ge Yang, Bradley C. Stadie, Roberto Calandra, Pieter Abbeel, Sergey Levine, Chelsea Finn

In Neural Information Processing Systems (NIPS) Deep RL workshop [Spotlight], Montreal, Canada, December 2018

Preprint. Available at: <https://github.com/bstadie/All-Bradly-Stadie-Papers/>

11. *One-Demonstration Imitation Learning*

Bradley C. Stadie, Siyan Zhao, Qiqi Xu, Bonnie Li, Lunjun Zhang

Submitted to ICLR 2020

Preprint. Available at: <https://github.com/bstadie/All-Bradly-Stadie-Papers/>

Invited Talks

- **Northwestern University Han Liu Group invited speaker (October 2019):** Learning reward functions.
- **INFORMS panel on ML in causal inference (October 2019):** Transfer learning causal effects.
- **University of Illinois, Urbana Champaign (October 2019):** Learning reward functions.
- **Toyota Technological Institute, Chicago (May 2019):** Learning from imperfect data.
- **University of Wisconsin, Madison (April 2019):** Learning from imperfect data.
- **Purdue University (April 2019):** Learning from imperfect data.
- **University of Montreal (February 2018):** The relationship between sampling and meta learning.
- **University of Toronto (February 2018):** The relationship between sampling and meta learning.
- **Harvard University 2017 Big Data Conference (August 2017):** Frontiers of meta learning.
- **Open AI (March 2016):** Exploration in deep reinforcement learning.

Teaching

- **Statistics 133: Concepts in Computing with Data.** Graduate Student Instructor. *Jan - May 2015*
- **Statistics 088: Introduction to Data Science.** Graduate Student Instructor. *Jan - May 2016*